

Assignment 5: Epipolar geometry and triangulation

Machine perception

2022/2023

Create a folder `assignment5` that you will use during this assignment. Unpack the content of the `assignment5.zip` that you can download from the course webpage to the folder. Save the solutions for the assignments as Python scripts to `assignment5` folder. In order to complete the assignment you have to present your solutions to the teaching assistant. Some assignments contain questions that require sketching, writing or manual calculation. Write these answers down and bring them to the assignment defense as well.

The maximum total amount of points for each assignment is 100. **All** the tasks in the assignment are mandatory, except if they are marked with \star . You can obtain 75 points for solving all the mandatory tasks. The remaining 25 points can be gained by solving optional tasks. There might be more than 25 points for optional assignments. In that case, the maximum of 100 points can **not** be exceeded, even if you solve all of them. The number of points for each optional task is written next to the instructions.

At the defense, the obligatory tasks **must** be implemented correctly and completely. If your implementation is incomplete, you will not be able to successfully defend the assignment.

Introduction

In this assignment we will review several parts of the epipolar geometry [1] (chapter 10.1) and robust estimation [1] (page 346).

Note: You can use feature point detection and matching that you implemented in Assignment 4. If you are unsure of your implementation or if you want to experiment with different detection and matching methods, you are allowed to use algorithms implemented in OpenCV, such as SIFT or ORB. Using functions such as `cv2.findFundamentalMat()` or `cv2.triangulatePoints()` is, of course, not allowed.

Exercise 1: Disparity

In this assignment we will focus on calculating disparity from a two-camera system. Our analysis will be based on a simplified stereo system, where two identical cameras are aligned with parallel optical axes and their image planes (CCD sensors) lie on the same plane (Image 1a).




Figure 1: Left image shows the simplest stereo setup with two parallel cameras. Right image shows geometric relations when mapping of a point \mathbf{p} in 3D space to axis x in image planes of the cameras.

In Figure 1b we can observe that we can write the following relations using similar triangles:

$$\frac{x_1}{f} = \frac{p_x}{p_z} \quad , \quad \frac{-x_2}{f} = \frac{T - p_x}{p_z}. \quad (1)$$

- (a) Using the equations (1) derive the expression for *disparity* which is defined as $d = x_1 - x_2$. What is the relation between the distance of the object (point \mathbf{p} in Figure 1) to camera and the disparity d ? What happens to disparity when the object is close to the cameras and what when it is far away?
 - (b) Write a script that computes the disparity for a range of values of p_z . Plot the values to a figure and set the appropriate units to axes. Use the following parameters of the system: focal length is $f = 2.5\text{mm}$ and stereo system baseline is $T = 12\text{cm}$.
 - (c) In order to get a better grasp on the idea of distance and disparity, you will calculate the numbers for a specific case. We will take the parameters from a specification of a commercial stereo camera *Bumblebee2* manufactured by the company PointGray: $f = 2.5\text{mm}$, $T = 12\text{cm}$, whose image sensor has a resolution of 648×488 pixels that are square and the width of a pixel is $7.4\mu\text{m}$. We assume that there is no empty space between pixels and that both cameras are completely parallel and equal. Lets say that we use this system to observe a (point) object that is detected at pixel 550 in x axis in the left camera and at the pixel 300 in the right camera. How far is the object (in meters) in this case? How far is the object if the object is detected at pixel 540 in the right camera? Solve this task analytically and bring your solution to the presentation of the exercise.
 - (d) ★ (10 points) Write a script that calculates the disparity for an image pair. Use the images in the directory **disparity**. Since the images were pre-processed we can limit the search for the most similar pixel to the same row in the other image. Since just the image intensity carries too little information, we will instead compare small

image patches. A simple way of finding a matching patch is to use normalized cross correlation. NCC for matrices \mathbf{X} and \mathbf{Y} of equal size is defined as

$$NCC(\mathbf{X}, \mathbf{Y}) = \frac{\sum(x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum(x_i - \bar{x})^2 \sum(y_i - \bar{y})^2}}. \quad (2)$$

where x_i denotes a specific cell in matrix \mathbf{X} and y_i a specific cell in matrix \mathbf{Y} . A patch from the second image is considered a match if it has the highest NCC value. The difference in x axis between the point from the first image and the matching point from the second image is the disparity estimate in terms of pixels¹. The disparity search is performed in a single direction.

Question: Is the disparity estimated well for all pixels? If not, what are the characteristics of pixels with more consistent disparity estimation?




Figure 2: Disparity calculated with patch size 10 and search window 35. Image was reduced to 50 % of its size.

- (e) ★ (5 points) Improve the noisy disparity estimation. You can implement a symmetric matching technique where the left image is compared to the right image and vice versa. The result can then be merged in some way that you choose. Smooth the results using median filter and choose an appropriate search window to obtain good results. Very noisy results will not be graded.

Note: You cannot merge the results element-wise. The images were taken from different viewpoints and the results would be incorrect.

Exercise 2: Fundamental matrix, epipoles, epipolar lines

In the previous exercise we were dealing with a special stereo setup where image planes of two cameras were aligned. In that case the projection of a 3D point has the same y coordinate in both cameras and different x coordinate. Such a setup ensures that the *epipolar lines in the cameras are parallel to the lines in the sensor*. This simplifies the search for correspondences (i.e. the projection of the same 3D point to both cameras). Generally the epipolar lines are not aligned with rows and in order to establish the relation between two image planes, the *fundamental matrix* needs to be computed.

¹Converting this estimate to distance requires information about the camera, which we do not have for the given image pairs.

In this exercise you will implement a simple version of an *eight-point algorithm* [1] that can be used to estimate a fundamental matrix between two cameras. We will first revisit the theory that will be used for the task.

We are given a list of perfect correspondence pairs of points in left $\mathbf{x} = [u, v, 1]^T$ and right $\mathbf{x}' = [u', v', 1]^T$ image in homogeneous coordinates (u corresponds to column, v to row). The fundamental matrix rule states that each correspondence is a valid solution of the following equation:

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0 \quad ; \quad \mathbf{F} = \begin{bmatrix} F_{11} & F_{12} & F_{13} \\ F_{21} & F_{22} & F_{23} \\ F_{31} & F_{32} & F_{33} \end{bmatrix}, \quad (3)$$

where \mathbf{F} denotes a fundamental matrix. Similarly to what we have done for the estimation of homography in the previous exercise, we can write a relation between a single pair of correspondence points as

$$\begin{bmatrix} F_{11} \\ F_{12} \\ F_{13} \\ F_{21} \\ F_{22} \\ F_{23} \\ F_{31} \\ F_{32} \\ F_{33} \end{bmatrix} = 0. \quad (4)$$

$$\begin{bmatrix} uu' & u'v & u' & uv' & vv' & v' & u & v & 1 \end{bmatrix}$$

If we combine $N \geq 8$ of these equations in a matrix \mathbf{A} we get a system of linear equations:

$$\begin{bmatrix} u_1 u'_1 & u'_1 v_1 & u'_1 & u_1 v'_1 & v_1 v'_1 & v'_1 & u_1 & v_1 & 1 \\ u_2 u'_2 & u'_2 v_2 & u'_2 & u_2 v'_2 & v_2 v'_2 & v'_2 & u_2 & v_2 & 1 \\ \vdots & \vdots \\ u_N u'_N & u'_N v_N & u'_N & u_N v'_N & v_N v'_N & v'_N & u_N & v_N & 1 \end{bmatrix} \begin{bmatrix} F_{11} \\ F_{12} \\ F_{13} \\ F_{21} \\ F_{22} \\ F_{23} \\ F_{31} \\ F_{32} \\ F_{33} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}. \quad (5)$$

We can solve the linear system above in a least-squares way by using the singular value decomposition method (SVD). Matrix \mathbf{A} is decomposed to $\mathbf{A} = \mathbf{U} \mathbf{D} \mathbf{V}^T$. The solution according to least squares corresponds to the eigenvector \mathbf{v}_n with the lowest eigenvalue, e.t. the last column of the matrix² \mathbf{V} .

Recall that the *epipole* of a camera is a point where all the epipolar lines (for that camera) intersect. This requires the rank of the matrix \mathbf{F} to be 2, however, this is usually not true when dealing with noisy data – the fundamental matrix estimated using the approach above will not have a rank 2. In practice this means that the epipolar lines will not intersect in a single point but rather in a small neighborhood of such a point. To stabilize the system we have to subsequently limit the rank of the fundamental matrix. This can be done by performing a SVD decomposition to $\mathbf{F} = \mathbf{U} \mathbf{D} \mathbf{V}^T$, set the lowest eigenvalue (D_{33} in matrix \mathbf{D}) to 0, and then reconstruct back the corrected matrix \mathbf{F} by

²This is a solution in column form as seen in equation (4) that has to be reshaped to the form in equation (3).

multiplying \mathbf{UDV}^T . A newly created fundamental matrix will satisfy the rank condition $\text{rank} = 2$ and can be used to compute two epipoles (one for each camera)

$$\mathbf{Fe} = 0 \quad \text{in} \quad \mathbf{F}^T \mathbf{e}' = 0, \quad (6)$$

by decomposing the matrix \mathbf{F} again and computing the *left* and *right*³ eigenvector of the matrix \mathbf{F} ,

$$\mathbf{e} = [V_{13} \ V_{23} \ V_{33}] / V_{33}, \mathbf{e}' = [U_{13} \ U_{23} \ U_{33}] / U_{33}, \quad (7)$$

which are then used to obtain both epipoles in homogeneous coordinates.

A simple eight-point algorithm for estimation of a fundamental matrix and the epipoles can be summarized as:

- Construct the matrix \mathbf{A} as in equation (5)
- Decompose the matrix using SVD $\mathbf{A} = \mathbf{UDV}^T$, and transform the last eigenvector \mathbf{v}_9 in a 3×3 fundamental matrix \mathbf{F}
- Decompose $\mathbf{F} = \mathbf{UDV}^T$ and set the lowest eigenvalue to 0, reconstruct $\mathbf{F} = \mathbf{UDV}^T$.
- Decompose $\mathbf{F} = \mathbf{UDV}^T$ again, compute both epipoles following the equation (7).

Once we have the fundamental matrix, we can take any point \mathbf{x} in the first image plane and determine an epipolar line for that point in the second image plane $\mathbf{l}' = \mathbf{Fx}$. Likewise, we can take point \mathbf{x}' in the second image plane and find the epipolar line in the first image plane $\mathbf{l} = \mathbf{F}^T \mathbf{x}'$. The line in the second image is then represented by all solutions of u' and v' in the equation $[u', v', 1]^T \mathbf{l}' = 0$.

- (a) Solve the following task analytically. We are given a system of two cameras and a fundamental matrix that connects the left camera to the right one

$$\mathbf{F} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0.5 & 0 \\ 0 & 0 & -1 \end{bmatrix}. \quad (8)$$

Compute the equation of the epipolar line in the right camera that corresponds to the point at column = 0 and row = 2 in the left camera. Take into account that the point has to be first written in homogeneous coordinates, i.e. $\mathbf{x} = [\text{column}, \text{row}, 1]^T$. Also compute the epipolar line for another point at column = 1 and row = 0 in the left camera, and the epipole.

Estimating a fundamental matrix

- (b) Implement a function `fundamental_matrix` that is given a set of (at least) eight pairs of points from two images and computes the fundamental matrix using the eight-point algorithm.

As the eight-point algorithm can be numerically unstable, it is usually not executed directly on given pairs of points. Instead, the input is first normalized by centering

³Attention: the terms left and right eigenvector are mathematical terms and do not hold any relation to the left and right camera in our system.

them to their centroid and scaling their positions so that the average distance to the centroid is $\sqrt{2}$. To achieve this, you can use the function `normalize_points` from the supplementary material.

Extend the function `fundamental_matrix` so that it first *normalizes* the input point-set of the left camera (we get transformed points and the transformation matrix \mathbf{T}_1) and then transform the input point set of the right camera (we get the transformed points and the transformation matrix \mathbf{T}_2). Using the transformed points the algorithm computes fundamental matrix $\hat{\mathbf{F}}$, then transforms it into the original space using both transformation matrices $\mathbf{F} = \mathbf{T}_2^T \hat{\mathbf{F}} \mathbf{T}_1$.

Test your function for fundamental matrix estimation using ten correspondence pairs that you load from the file `house_points.txt`. The columns are formatted as follows: x_1, y_1, x_2, y_2 , i.e. the first column contains the x -coordinates of the points for the first image etc. Compute the fundamental matrix \mathbf{F} and for each point in each image calculate the corresponding epipolar line in the other image. You can draw the epipolar lines using `draw_epiline` from the supplementary material. According to epipolar geometry the corresponding epipolar line should pass through the point. As a testing reference the correct fundamental matrix is included in the supplementary material in file `house_fundamental.txt`.




Figure 3: Correspondences and epipolar lines for both images, calculated with the fundamental matrix.

- (c) We use the *reprojection error* as a quantitative measure of the quality of the estimated fundamental matrix.

Write a function `reprojection_error` that calculates the reprojection error of a fundamental matrix F given two matching points. For each point, the function should calculate the corresponding epipolar line from the point's match in the other image, then calculate the perpendicular distance between the point and the line using the equation:

$$distance(ax + by + c = 0, (x_0, y_0)) = \frac{|ax_0 + by_0 + c|}{\sqrt{a^2 + b^2}}, \quad (9)$$

where a , b and c are the parameters of the epipolar line. Finally, the function should return the average of the two distances.

Write a script that performs two tests: (1) compute the reprojection error for points $p_1 = [85, 233]^T$ in the left image-plane and $p_2 = [67, 219]^T$ in right image-plane using the fundamental matrix (the error should be approximately 0.15 pixels). (2) Load

the points from the file `house_points.txt` and compute the average of symmetric reprojection errors for all pairs of points. If your calculation is correct, the average error should be approximately 0.33 pixels.

- (d) ★ (15 points) Perform fully automatic fundamental matrix estimation on a pair of images from the directory `desk`⁴. Detect the correspondence points using your preferred method. As some of the matches might be incorrect, extend the RANSAC algorithm so that it will work for fundamental matrix estimation. You can measure the quality of a solution by using the point-to-line reprojection error. Display the correspondences to check whether all of them are correct. Calculate the fundamental matrix on the final set of inliers and show correspondences with epipolar lines.

Exercise 3: Triangulation

If we know the intrinsic parameters of the cameras, we can calculate the 3D position of the points observed in both cameras. You can find the projection matrices for both cameras stored in files `house1_camera.txt` and `house2_camera.txt`. The calibration matrices have the size 3×4 .

We will use an algebraic approach to perform the triangulation. Assuming we have a 2D correspondence between \mathbf{x}_1 in the first image plane and \mathbf{x}_2 in the second image plane (in homogeneous coordinates), a location of the common point \mathbf{X} in 3D space (4D in homogenous coordinates) is given by relations

$$\lambda_1 \mathbf{x}_1 = \mathbf{P}_1 \mathbf{X} \quad (10)$$

$$\lambda_2 \mathbf{x}_2 = \mathbf{P}_2 \mathbf{X}. \quad (11)$$

where $\lambda_1, \lambda_2 \in \mathbb{R}$. We know that a vector product between parallel vectors is 0 so we use $\mathbf{x}_1 \times \lambda_1 \mathbf{x}_1 = 0$ and get:

$$\mathbf{x}_1 \times \mathbf{P}_1 \mathbf{X} = [\mathbf{x}_{1\times}] \mathbf{P}_1 \mathbf{X} = 0 \quad (12)$$

$$\mathbf{x}_2 \times \mathbf{P}_2 \mathbf{X} = [\mathbf{x}_{2\times}] \mathbf{P}_2 \mathbf{X} = 0, \quad (13)$$

where we have used the following form (shear-symmetric form) to get rid of a vector product:

$$\mathbf{a} \times \mathbf{b} = [\mathbf{x}_{1\times}] \mathbf{b} = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ -a_y & a_x & 0 \end{bmatrix} \mathbf{b}. \quad (14)$$

For each pair of 2D points we get two independent linear equations for three unknown variables. If we combine in a matrix \mathbf{A} the first two lines of the product $[\mathbf{x}_{1\times}] \mathbf{P}_1$ and first two lines of the product $[\mathbf{x}_{2\times}] \mathbf{P}_2$, we can compute the mean quadratic estimate of \mathbf{X} by solving a linear equation system $\mathbf{AX} = 0$. As you already know by now, such a solution can be obtained by computing the eigenvector of matrix \mathbf{A} that has the lowest eigenvalue. Note that the solution of the system is a point \mathbf{X} in homogeneous coordinates (4D space), therefore you have to first normalize the values so the last coordinate becomes 1.

⁴Credit: <http://nghiaho.com/>

- (a) Implement the function `triangulate` that accepts a set of correspondence points and a pair of calibration matrices as an input and returns the triangulated 3D points. Test the triangulation on the ten points from the file `house_points.txt`. Visualize the result using `plt.plot` or `plt.scatter`. Also plot the index of the point in 3D space (use `plt.text`) so the results will be easier to interpret. Plot the points interactively, so that you can rotate the visualization.

Note: The coordinate system used for plotting in 3D space is usually not the same as the camera coordinate system. In order to make the results easier to interpret, the ordering of the axes can be modified by using a transformation matrix. For example, you may want to use matrix \mathbf{T} :

$$\mathbf{T} = \begin{bmatrix} -1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & -1 & 0 \end{bmatrix}. \quad (15)$$

to transform 3D points with $\mathbf{T}\mathbf{x}$.




Figure 4: Points on the first and second image, and their 3D positions visualized in a plot. To get a feeling of their placement, rotate the plot.

- (b) ★ (25 points) Perform a 3D reconstruction of an object you own. For that you will need to calibrate a camera with which you will take images of the object. You can use a webcam or a cell phone. Print out a calibration pattern⁵ and take multiple images with the pattern visible on a planar surface (if you scale it correctly you can also take pictures of your screen). Take care that you change the orientation and distance between the pattern and the camera. Detect the circle centers using `cv2.findCirclesGrid` and check the correctness with `cv2.drawChessboardCorners`. Finally, use `cv2.calibrateCamera` to obtain the camera intrinsic parameters from the detected patterns. You can check that applying `cv.undistort()` removes the potential distortion from your images (depending on the lens used, the changes might be small). If you performed the calibration correctly, the returned reprojection error should be below 1 pixel.

When you have the intrinsic parameters of your camera, you will need to take two images of your object from different viewpoints. First, undistort your images using `cv2.undistort`, then detect and match feature points. Calculate the fundamental matrix with RANSAC and, using the calibration matrix, also calculate the essential matrix. Then, you can use `cv2.recoverPose` to obtain the rotation and translation

⁵You can use this one: <https://nerian.com/nerian-content/downloads/calibration-patterns/pattern-a4.pdf>

parameters (extrinsics) for both camera viewpoints. Use the extrinsic parameters to calculate the projection matrices for both cameras and triangulate the matched points. Display the final 3d points. To further verify your solution, project the detected points to each of the images using the recovered projection matrices.

Note: `cv2.recoverPose` will only return one rotation matrix and one translation vector. This is because the first camera lies in the origin of the coordinate system and its rotation matrix equals to the 3×3 identity matrix, while its translation vector only contains zeros (naturally excluding the last element in homogenous coordinates).

References

- [1] D. A. Forsyth and J. Ponce. *Computer Vision: A Modern Approach*. Prentice Hall, 2002.